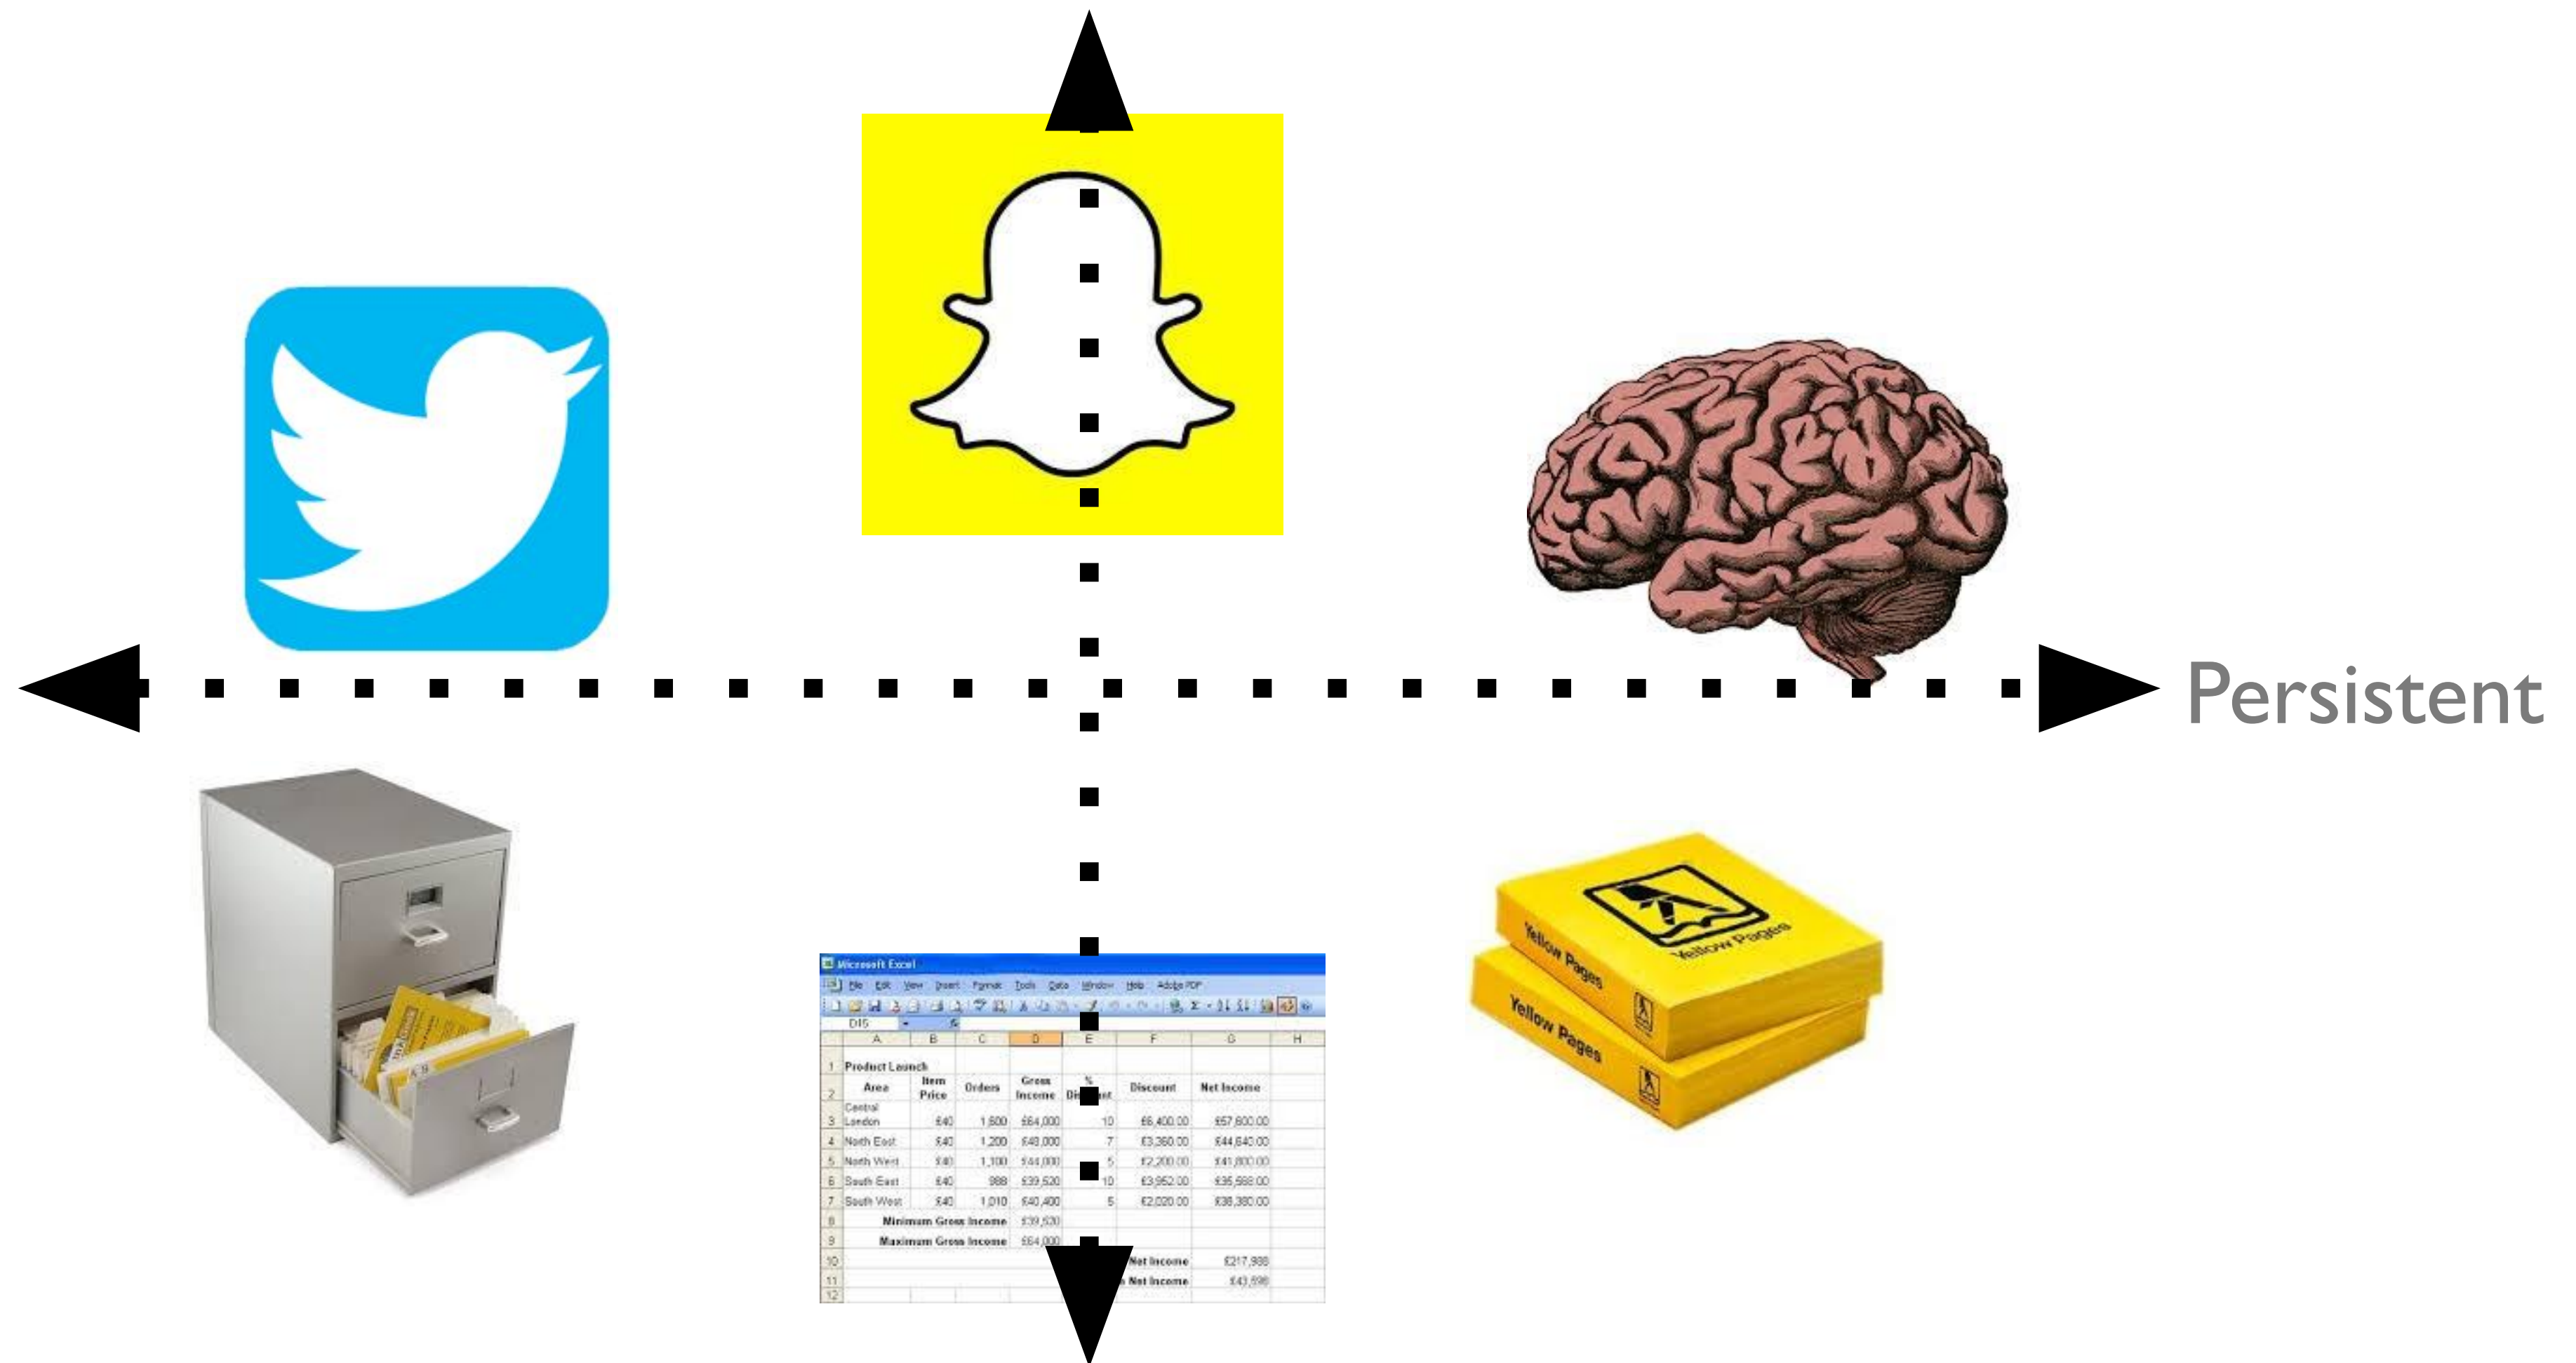


Intro to Databases

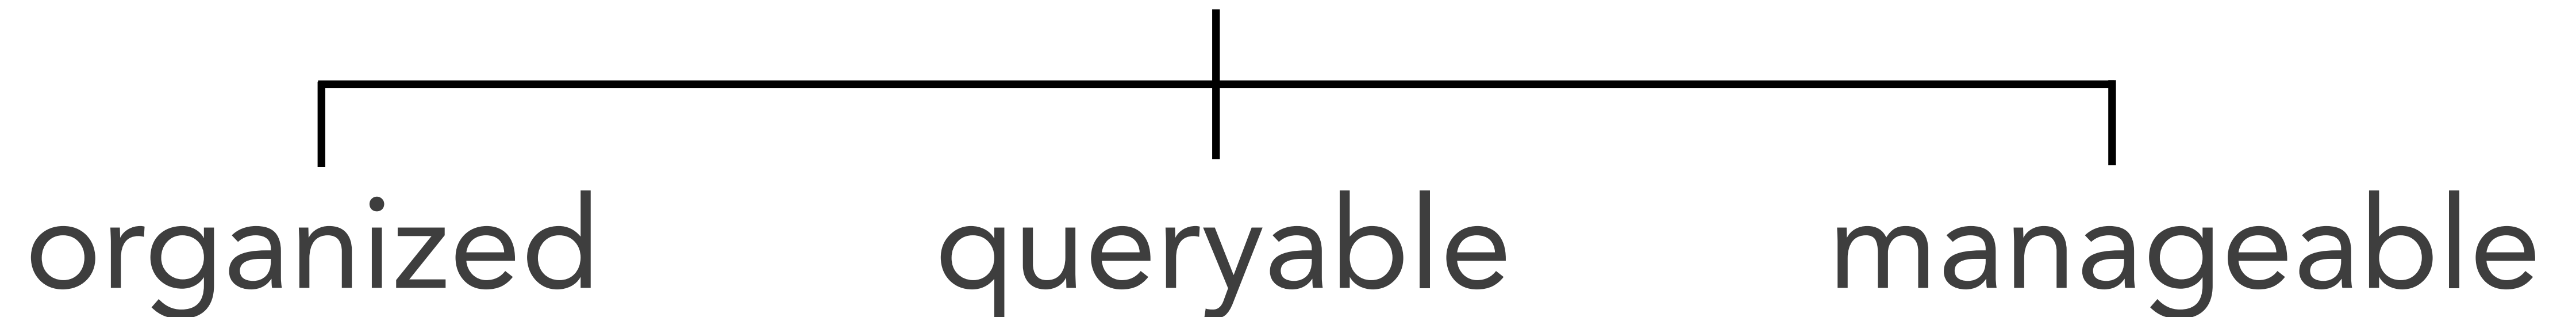
What is a database?

Things that hold info

Accessible



A database persists information and
is accessible via code



Organized: Standard Storage Formatting

- ◉ DBs are a collection of **Tables** (or *relations*)
- ◉ Tables have **Columns** (*attributes* / *fields*) that describe **Rows** (*instances* / *tuples*)
- ◉ Each row should be unique
- ◉ Rows often have a **primary key** (unique identifier)

Table / Relation

Column / Attribute / Field

Column / Attribute / Field

Column / Attribute / Field

ID	Name	Type
1	Pikachu	lightning
2	Squirtle	water
3	Charmander	fire
4	Bulbasaur	grass

Row / Tuple / Instance

Row / Tuple / Instance

Row / Tuple / Instance

Row / Tuple / Instance

Queryable: via a Standard Language

- ◉ A simple, structured query language: SQL
 - Declarative (vs. imperative)
 - No more hand-rolled algorithms / data structures
 - DBMS picks an efficient execution strategy based on indexes, data, workload etc.



SQL

```
-- Pikachu, I choose you!  
SELECT id, name  
FROM pokemon  
WHERE type = 'lightning'  
LIMIT 1
```


Manageable: Easy, Safe, Performant

- Offloads work and requisite understanding of programming
- Knowledge is portable
- Abstraction
- Transfer data between systems
- DBMS can make certain guarantees
 - prevent unsafe operations
 - built-in redundancies
 - handle multiple users, threads

ACID Guarantees

- ◉ Atomicity
- ◉ Consistency
- ◉ Isolation
- ◉ Durability

What happens if the database crashes?

- ◉ Suppose we have a banking system with client accounts
- ◉ Every debit must have a matching credit.
- ◉ Imagine a crash results in only one table being updated.
 - Database inconsistency → unexpected data and software crashes/bugs
 - Financial risk for clients

Atomic Transactions

- ◉ atomic transaction: A set of database operations that must occur together
 - i.e. A debit to one bank account, and a credit to another
- ◉ A transaction must either succeed or fail; it cannot partially complete.
- ◉ Every database query is represented by a transaction

Consistency

- ◉ Specify rules that columns need to follow. Examples:
 - Savings account must start with S or checking with C
 - Column cannot be null
- ◉ Protect the database from inconsistencies and simplify software logic
 - Allows software to make assumptions about underlying data

Resource Management

- ◉ Processes can be readers and writers
- ◉ Files can have many readers
- ◉ If a process has a writer, no other process can read from it, and no other process can write to it

Proposed File Scheme

- ◉ Suppose that we have decided not to use a database and instead store our data in a series of files.
- ◉ How might our setup fail to serve queries from multiple users?

Deadlock



Databases give us concurrency (Isolation)

- ◉ Multiple clients can make queries to read and update without the risk of deadlock or starvation.

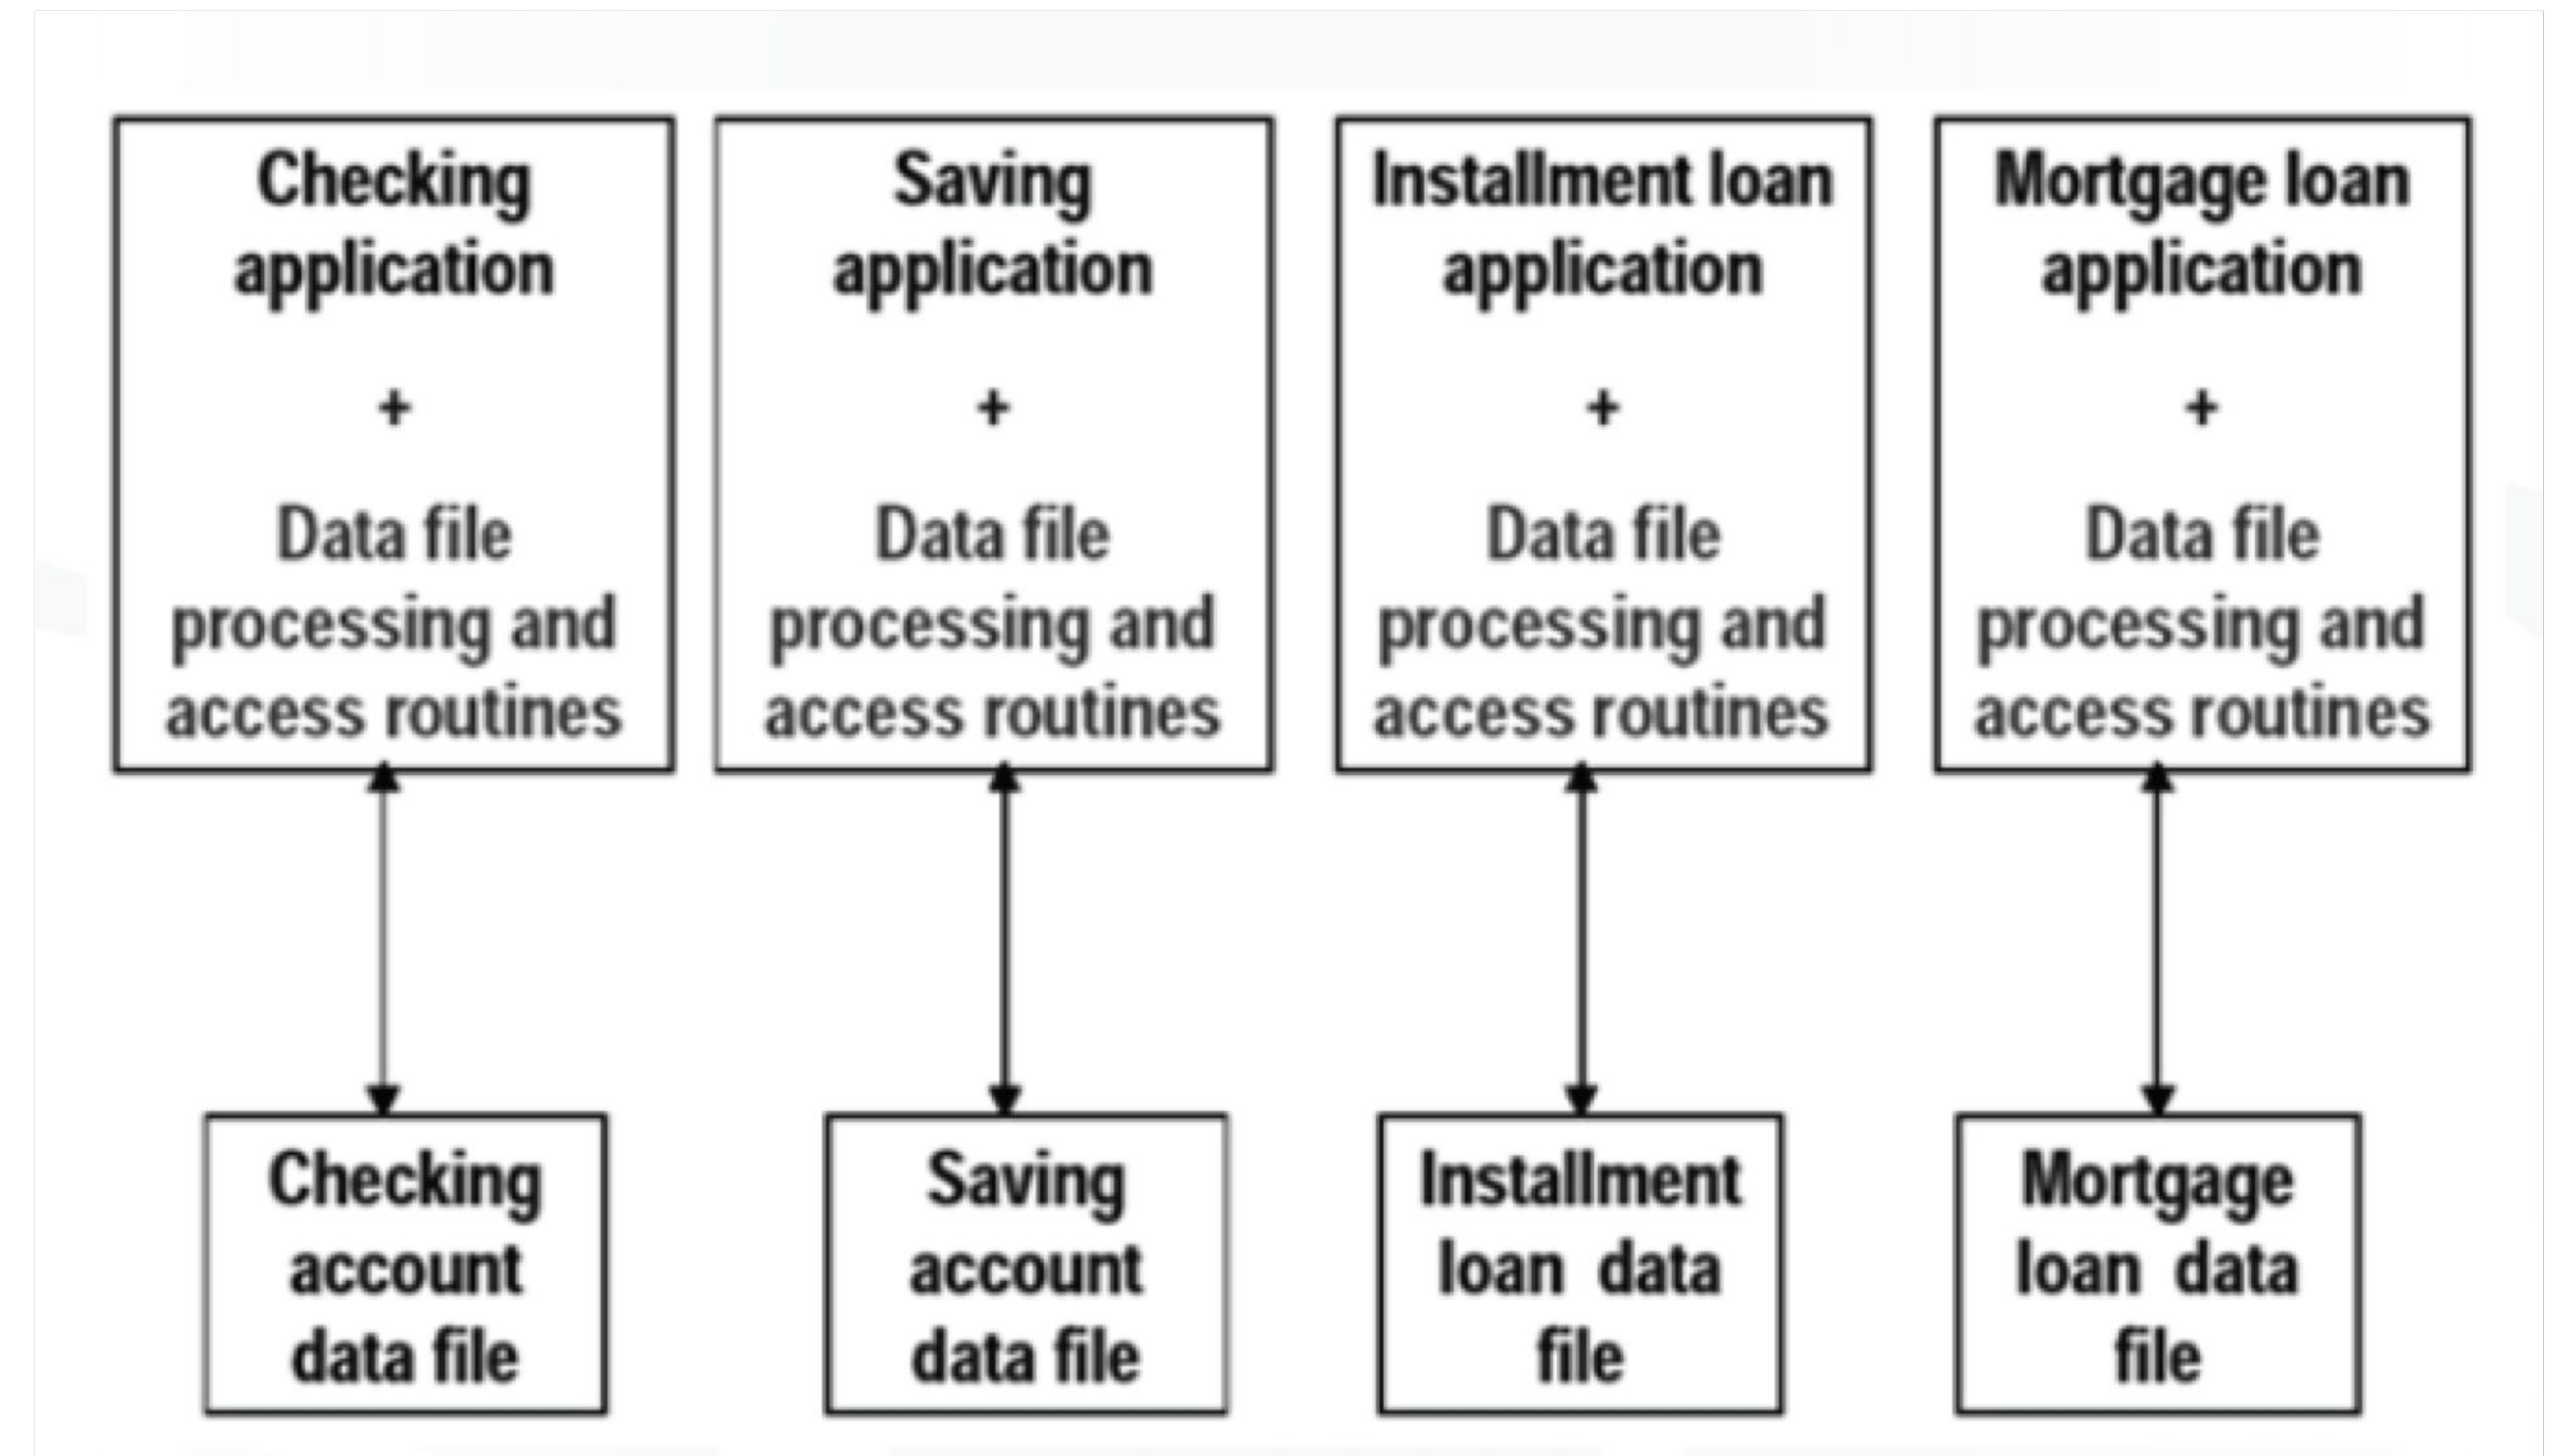
Persistence/Durability

- ◉ Databases are also persistent (store information without power)

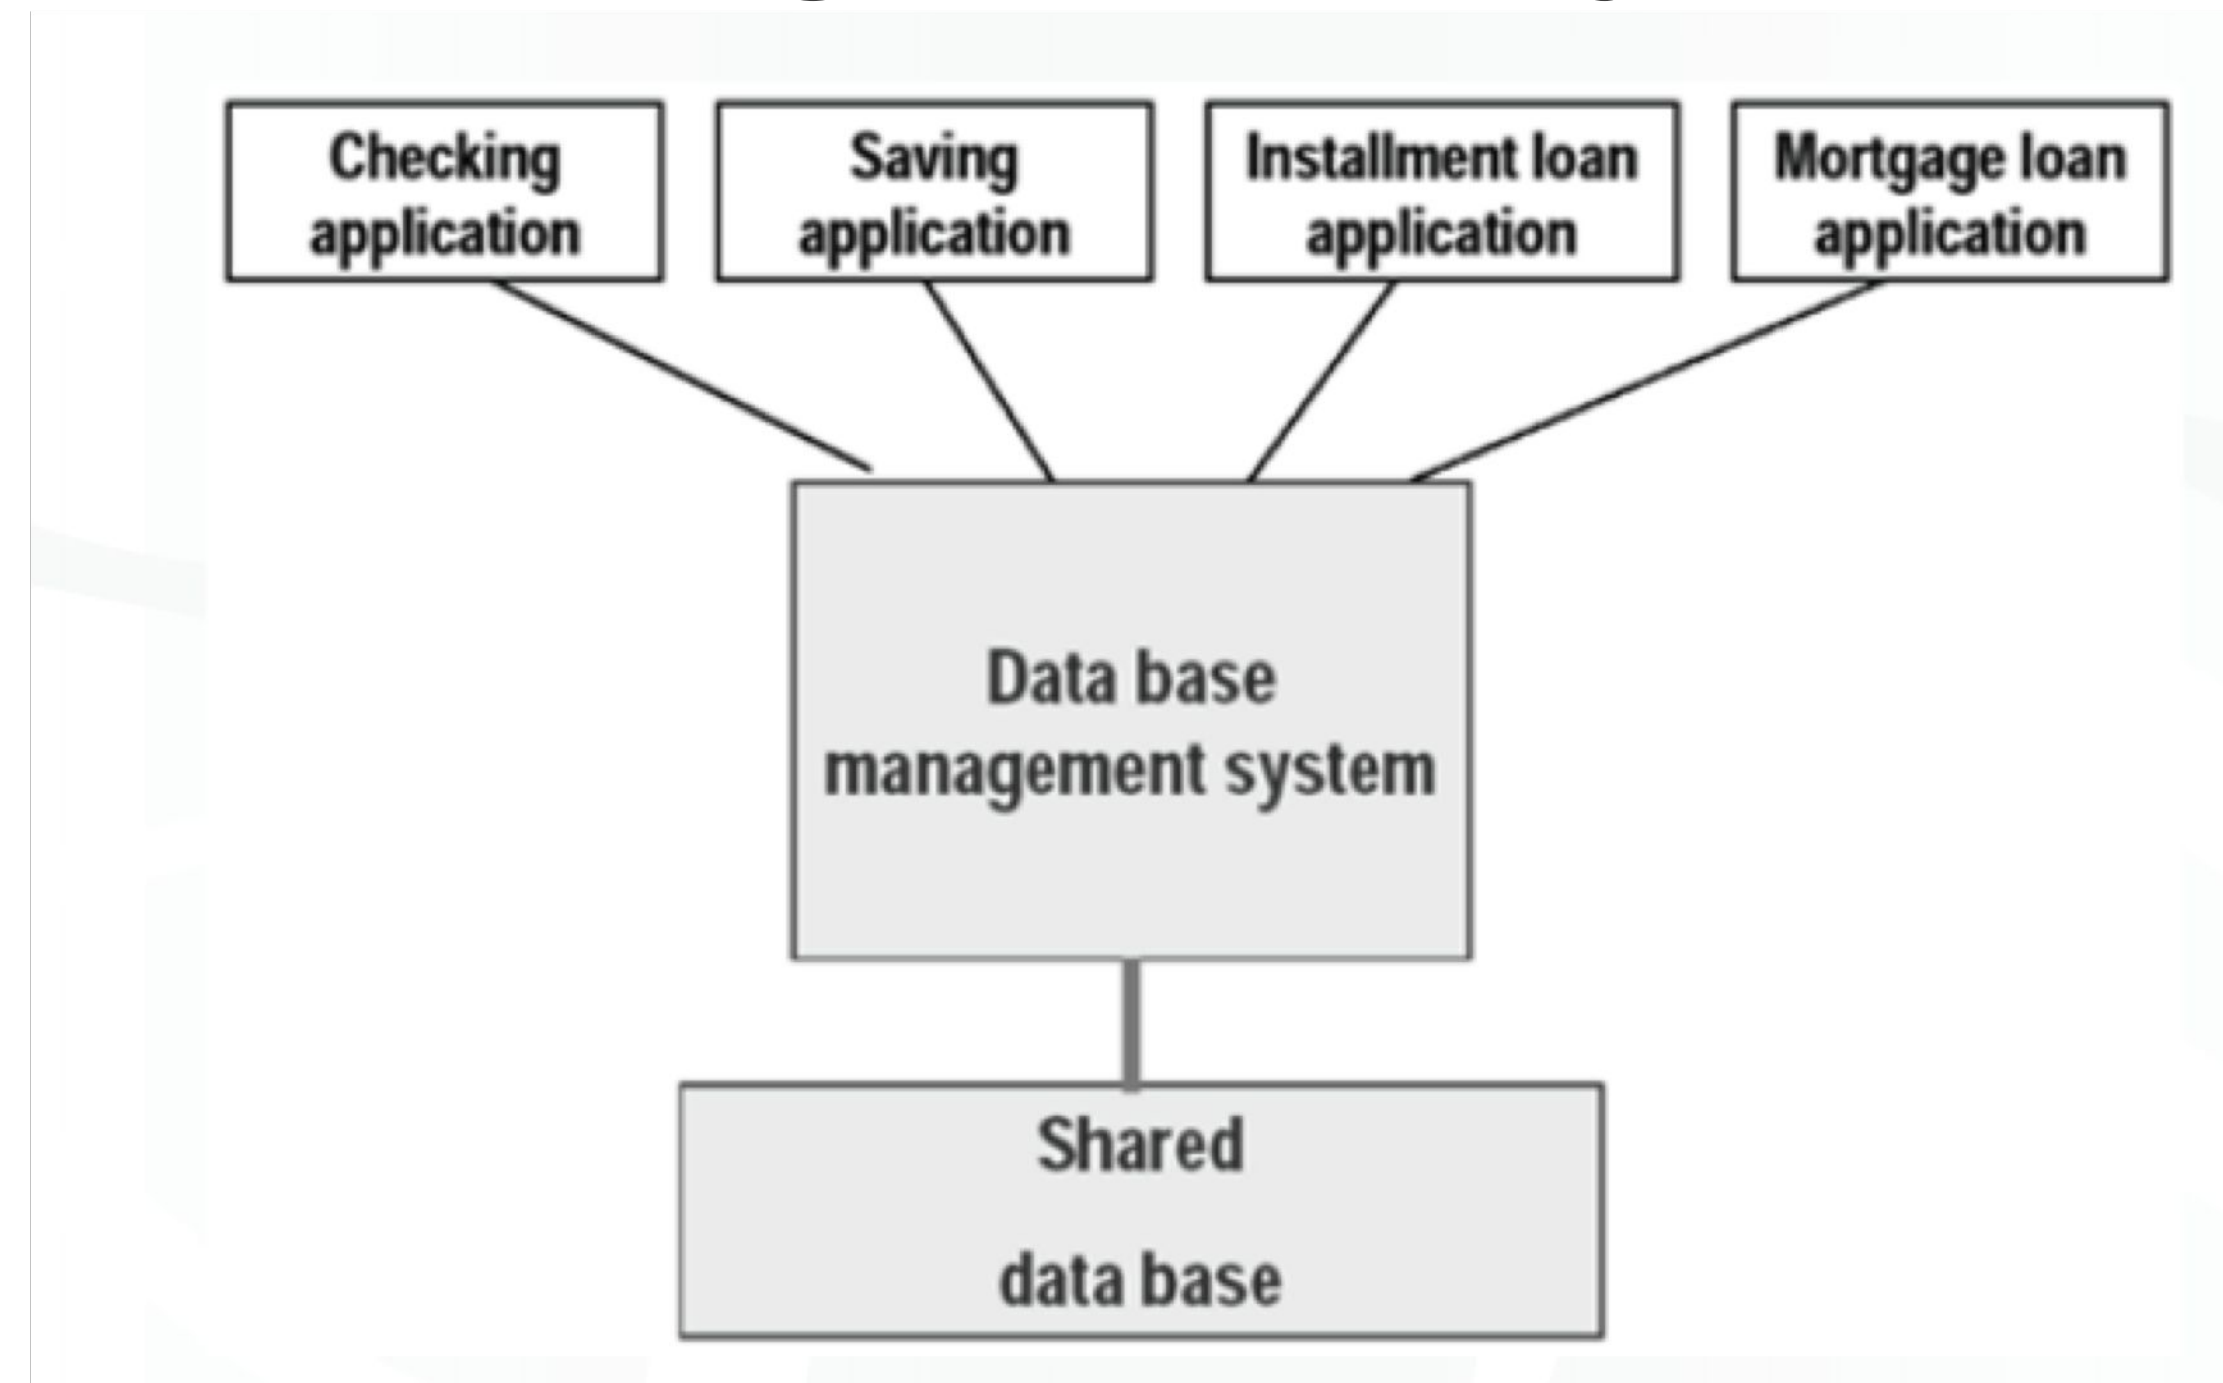
How Did We End Up Here?

Before Relational DBs (ca. < 1970s)

- ◉ Data stored in custom “data files”
- ◉ Queried via application-specific code
- ◉ Advantages
 - Middle layer not needed
 - Solutions customized for each application
- ◉ Disadvantages
 - Hard to change the system
 - Knowledge not compounding
 - Data-transfer is difficult



Database Management Systems (DBMS)

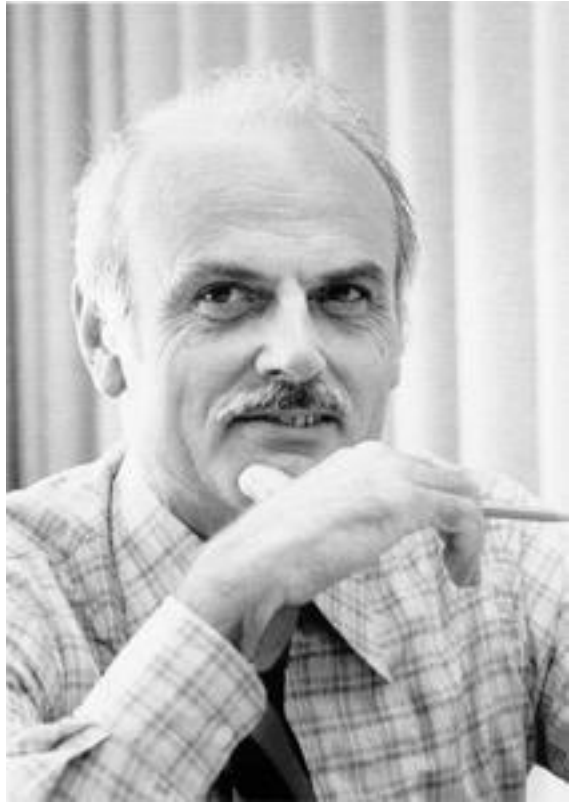


- ◉ One layer and language to store and access data
- ◉ Sold as a way for “non-technical people” to manage data

“Future users of large data banks must be protected from having to know how the data is organized in the machine (the internal representation).”

– E. F. CODD,
A RELATIONAL MODEL OF DATA FOR
LARGE SHARED DATA BANKS

Relational Databases & Logic



- ◉ 1969: Edgar Frank "Ted" Codd outlines *relational model* of data
- ◉ Wrote Alpha (never implemented) as a *query language*
- ◉ IBM slow to adopt his ideas
 - Competitors started to do so
 - IBM team formed without Codd, created **Structured English Query Lang**
- ◉ SEQUEL way better than what came before
 - 1979: copied by Larry Ellison (from pre-launch papers / talks!) as "SQL"
- ◉ SQL became the standard (ANSI 1986, ISO 1987)
 - Codd continued to fault SQL compared to his theoretical model
 - The Third Manifesto: solve the *object-relational impedance mismatch*

Appreciating Databases

- ◉ Ubiquitous
- ◉ Standardized
- ◉ Complex / deep
- ◉ Powerful: database admins are
 - Feared by developers
 - ...but also taken for granted until things break
 - Befriended by business people
 - Contacted by the government for secret data (e.g. NSA)

Progression of Databases

- ◉ Navigational (< 1970s)
 - More common during tape era; entries had references to next entries.
- ◉ Relational (> 1970s)
 - Based on relational (table-based) logic, see E.F. Codd.
- ◉ NoSQL (> 2000s)
 - "Not only SQL" — document storage, for example.

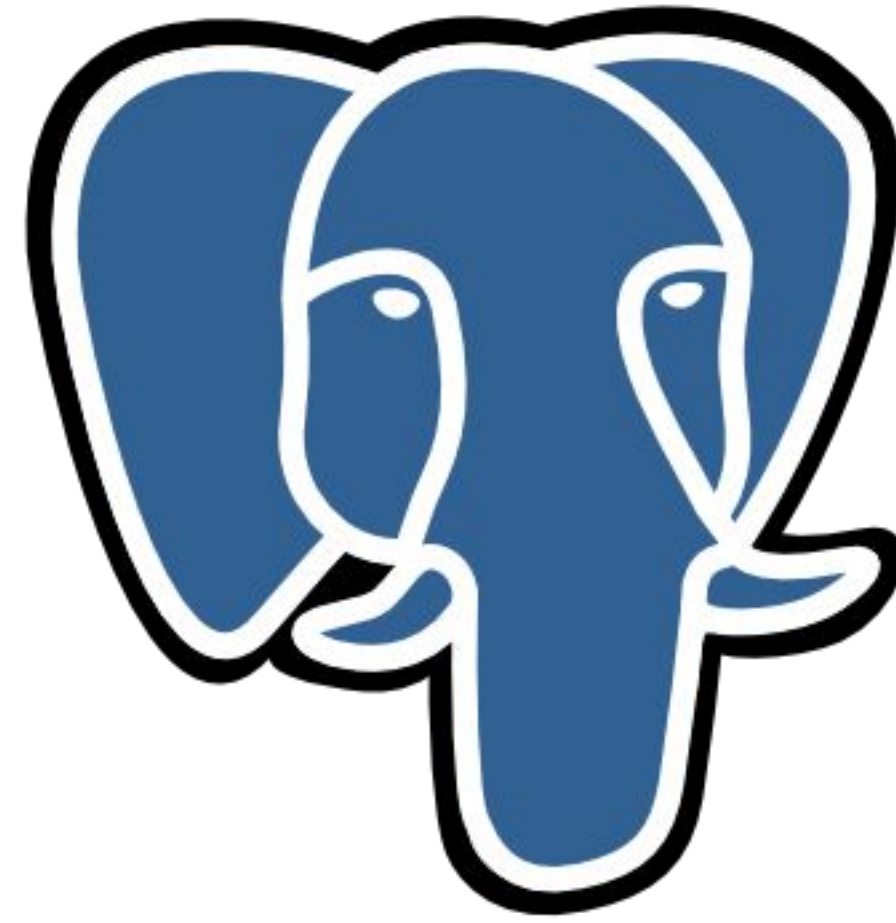
RDBMS vs NoSQL

- ◉ A DBMS doesn't ***have*** to be relational
 - Remember, DBMS is just an application that intelligently stores data and can answer requests to manage that data
- ◉ Lately, many "NoSQL" or non-relational DBMSs have been gaining popularity
 - Graph databases (e.g. Neo4J)
 - Document databases (e.g. MongoDB)
 - Hybrids (e.g. PostgreSQL)
- ◉ RDBMSs still remain the #1 DB option for now



Some well-known rDBMSs





PostgreSQL

Why PostgreSQL?

- ◉ Advanced, powerful, and popular
- ◉ Rapid open source development
- ◉ Highly extensible (stored procedures)
- ◉ Deep SQL standards compliance
- ◉ NoSQL ("Not Only SQL"), objective support
- ◉ Excellent transactions / ACID reliability; focus on integrity
- ◉ Multi-user management / administration

History of PostgreSQL

- ◉ 1970s at UC Berkeley:
Interactive **G**raphics **R**etrieval **S**ystem (INGRES)
- ◉ 1980s: POSTGRES ("Post-Ingres")
- ◉ 1995: POSTQUEL and Postgres95.
 - monitor -> psql
- ◉ 1996: Adopted by the open source community
 - Ongoing: stability, testing, documentation, new features
 - PostgreSQL

psql

```
1. psql (psql)
X .../hello-world (zsh)  %1 X ...s/class-libs (zsh)  %2 X psql (psql)  %3 X postgres (postgres)  %4
psql [local]:5432 glebec # ~ \l
List of databases

```

Name	Owner	Encoding	Collate	Ctype	Access privileges
assessmentexpresssequelize	fullstack	UTF8	en_US.UTF-8	en_US.UTF-8	
auther	glebec	UTF8	en_US.UTF-8	en_US.UTF-8	
checkpoint_angular	glebec	UTF8	en_US.UTF-8	en_US.UTF-8	
checkpoint_express_review	glebec	UTF8	en_US.UTF-8	en_US.UTF-8	
glebec	glebec	UTF8	en_US.UTF-8	en_US.UTF-8	
juke	glebec	UTF8	en_US.UTF-8	en_US.UTF-8	
postgres	glebec	UTF8	en_US.UTF-8	en_US.UTF-8	
sequelizecheckpoint	fullstack	UTF8	en_US.UTF-8	en_US.UTF-8	
template0	glebec	UTF8	en_US.UTF-8	en_US.UTF-8	=c/glebec ↵ glebec=CTc/glebec
template1	glebec	UTF8	en_US.UTF-8	en_US.UTF-8	=c/glebec ↵ glebec=CTc/glebec
tripplanner	glebec	UTF8	en_US.UTF-8	en_US.UTF-8	
twitterdb	glebec	UTF8	en_US.UTF-8	en_US.UTF-8	
wikistack	glebec	UTF8	en_US.UTF-8	en_US.UTF-8	

```

(13 rows)

psql [local]:5432 glebec # ~ \c auther
You are now connected to database "auther" as user "glebec".

psql [local]:5432 glebec # auther \dt
List of relations

```

Schema	Name	Type	Owner
public	stories	table	glebec
public	users	table	glebec

```

(2 rows)

psql [local]:5432 glebec # auther |
```



pgcli

```
stayupdated_test> \d
+-----+-----+-----+-----+
| Schema | Name                | Type      | Owner  |
+-----+-----+-----+-----+
| public | admins              | table     | amjith |
| public | cpes                | table     | amjith |
| public | goose_db_version    | table     | amjith |
| public | goose_db_version_id_seq | sequence  | amjith |
| public | packages            | table     | amjith |
| public | packages_id_seq     | sequence  | amjith |
| public | users               | table     | amjith |
| public | users_id_seq        | sequence  | amjith |
| public | vulnerabilities      | table     | amjith |
| public | vulnerabilities_cpes | table     | amjith |
| public | vulnerabilities_id_seq | sequence  | amjith |
+-----+-----+-----+-----+
SELECT 11
stayupdated_test> SELECT * FROM users;
+-----+-----+-----+-----+-----+
| id | display_name | password | email           | created_on |
+-----+-----+-----+-----+-----+
| 177 | DisplayName1 | 1024cms  | user@ex.com     | 2014-11-15 15:02:50.094560 |
| 180 | testname2    | pas5w0rd | email@ex.com    | 2014-11-28 10:25:46.170660 |
| 181 | amjith       | password | amjith@amjith.amjith | 2014-11-28 18:39:48.195067 |
+-----+-----+-----+-----+-----+
SELECT 3
stayupdated_test> SELECT * FROM
| admins
| cpes
| goose_db_version
| packages
| users
```


Postico

The screenshot shows the Postico application window. The top bar includes navigation icons, a breadcrumb trail (Reporting > reporting > forex), a status indicator (Connected.), the database version (PostgreSQL 9.4.5), and window controls. On the left is a sidebar with a list of database tables, including 'forex' which is currently selected. The main area displays a table with the following columns: currency, base, rate, date, and source_id. The 'forex' table is highlighted in blue. The table contains 16 rows of data, all for the date 2014-07-25. The 'rate' column shows various exchange rates against the base currency 'USD'. On the right side of the table, there are filters for 'currency' (set to 'MULTIPLE'), 'base' (set to 'USD'), 'rate' (set to 'MULTIPLE'), and 'date' (set to '2014-07-25'). Below these filters is a 'source_id' filter with a dropdown menu showing '1' and 'Open Exchange R...'. At the bottom of the window, there is a search bar, tabs for 'Content' and 'Structure', a '+' button, a 'Filter' button, and pagination information (Page 1 of 342).

currency	base	rate	date	source_id
AED	USD	3.67291	2014-07-25	1
AFN	USD	56.485726	2014-07-25	1
ALL	USD	103.5838	2014-07-25	1
AMD	USD	410.086	2014-07-25	1
ANG	USD	1.787	2014-07-25	1
AOA	USD	96.952626	2014-07-25	1
ARS	USD	8.169642	2014-07-25	1
AUD	USD	1.062844	2014-07-25	1
AWG	USD	1.79	2014-07-25	1
AZN	USD	0.784067	2014-07-25	1
BAM	USD	1.453024	2014-07-25	1
BBD	USD	2	2014-07-25	1
BDT	USD	77.61971	2014-07-25	1
BGN	USD	1.452543	2014-07-25	1

DEMO